



# **AI in Drug Discovery – An Overview**

## **Session 1**

**September 16, 2024**

# Who we are!

---

**Pat Walters**  
Relay Therapeutics



**Raquel López-Ríos de Castro**  
Charité Berlin, MSKCC NYC



**Michael Backenköhler**  
Saarland University



**Andrea Volkamer**  
Saarland University



# What we will do today

---

## Session 1 - 1:30 - 2:30 pm

- An introduction to Artificial Intelligence (AI) and Machine Learning (ML)
- Molecular representations
- AI architectures

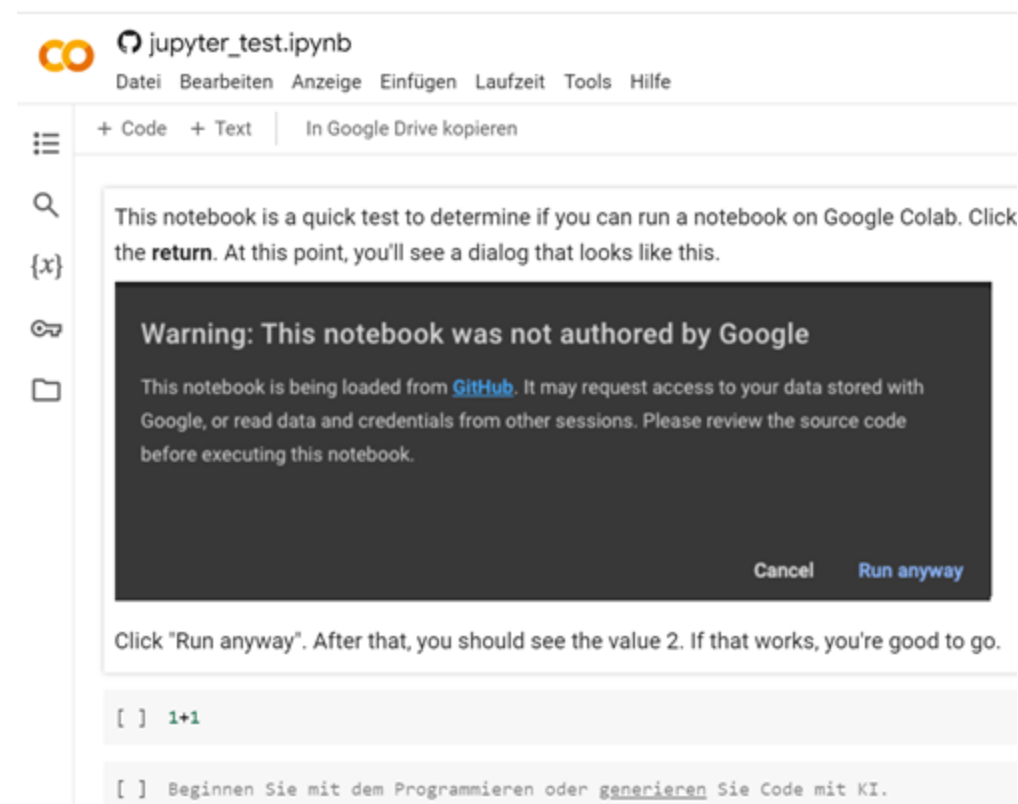
## Session 2 - 3:00 - 4:00 pm

- The importance of data quality for AI/ML
- Exploratory data analysis
- Data preprocessing
- Applicability domains

## Session 3 - 4:30 - 5:30 pm

- AI in Practice
- Molecule generation
- Active learning

Lectures supported by hands-on sessions ...



**artificial intelligence (AI)**, the ability of a digital **computer** or computer-controlled **robot** to perform tasks commonly associated with intelligent beings.

Not a well-defined statement



## AI and “The Rise of the Machines”

---

 Andrew Chen Retweeted



**Mat Velloso** @matvelloso · Nov 22

Difference between machine learning and **AI**:

If it is written in Python, it's probably machine learning

If it is written in **PowerPoint**, it's probably **AI**

 166

 6.6K

 19K



Show this thread

# What Is Machine Learning?

---

**Machine learning is all about labeling things using examples**




**Cassie Kozyrkov, Google**




# Labeling Molecules Based on Examples

Molecules with measured data




Molecules to be predicted



Log10(Molar Aqueous Solubility)


# Using Predictive Models to Drive Drug Discovery




On-target Activity



Off-target Activity



Pharmacokinetics

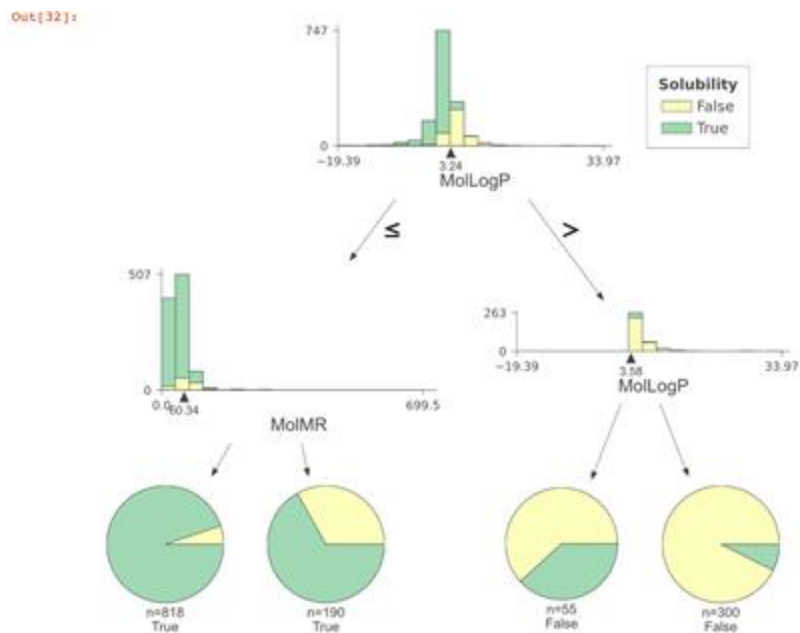


Physical Properties




# Two Types of ML Models – Classification and Regression

## Classification



Predict a category, e.g.  
soluble/insoluble


## Regression



Experimental  
LogS

Predict a value, e.g. 4.2

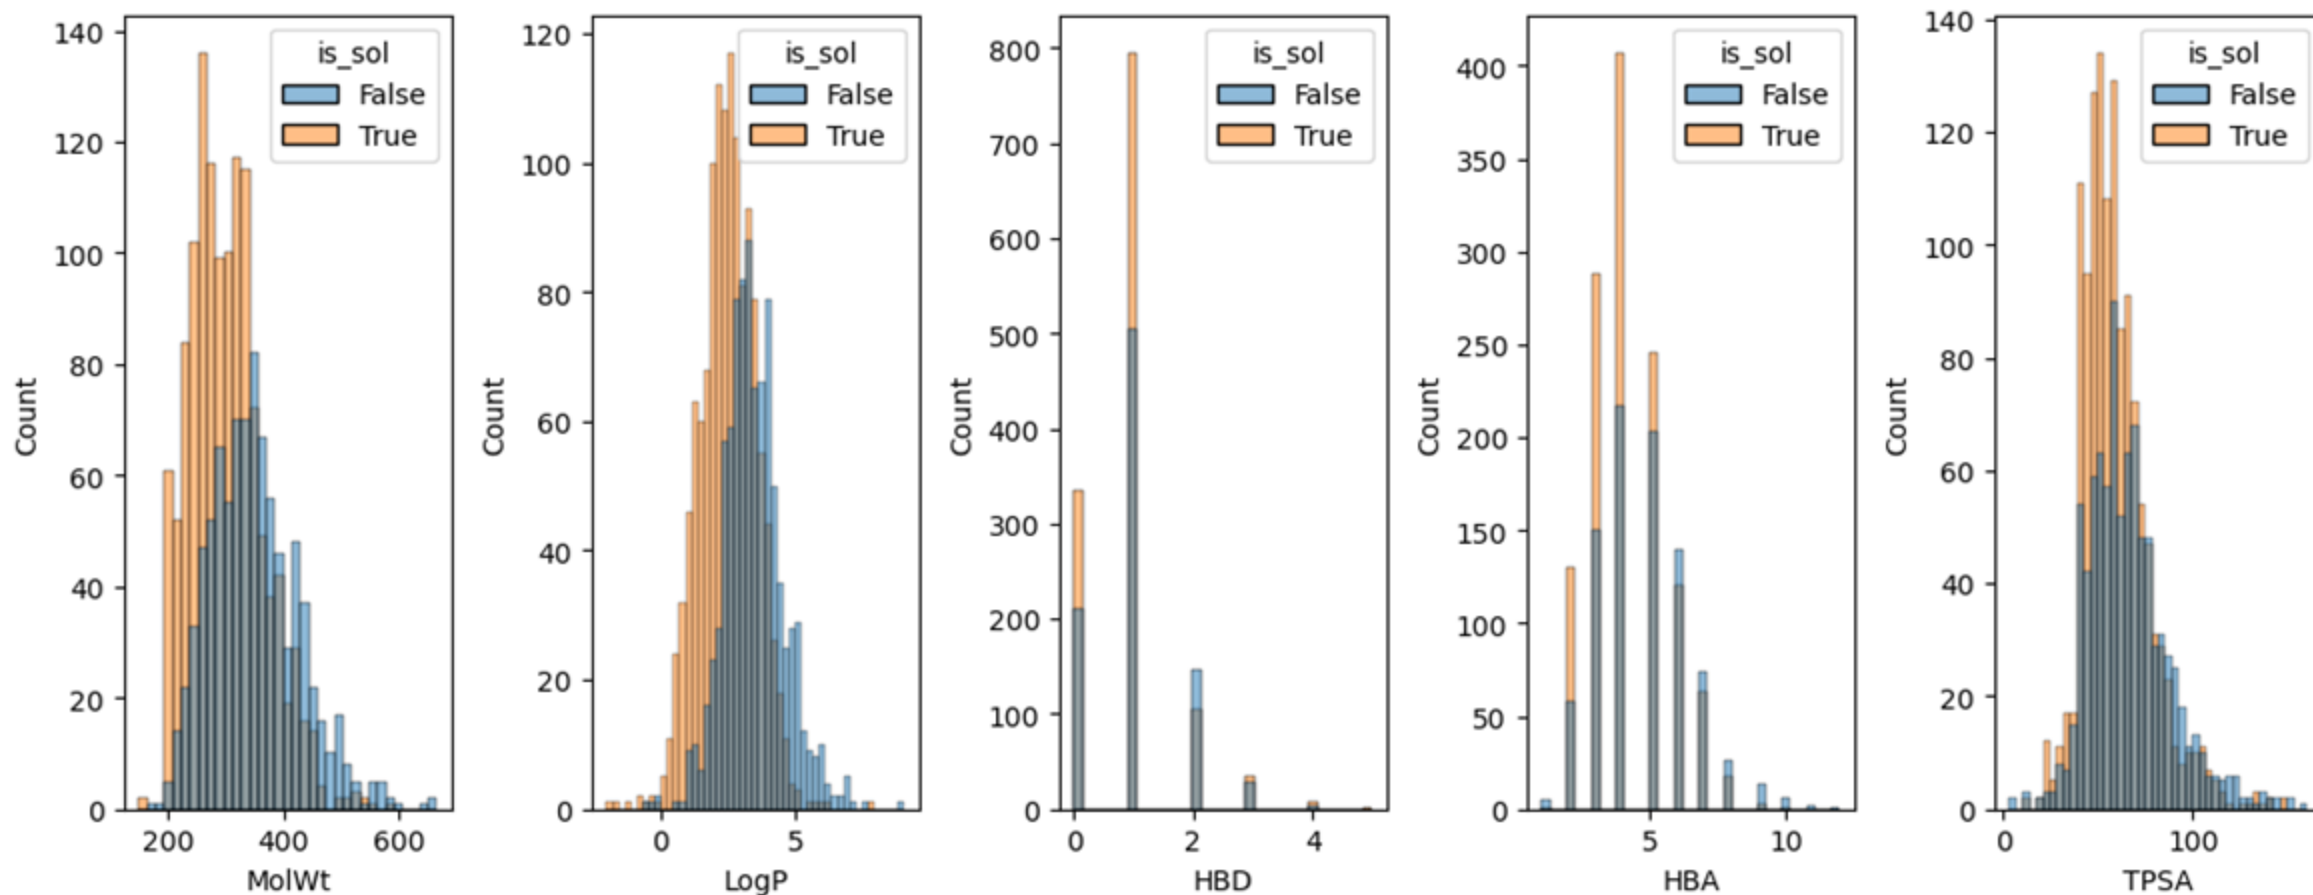
## Let's Start With a Dataset



Aqueous solubility by CLND  
2173 compounds  
Min = 0.23 $\mu$ M  
Max = 648 $\mu$ M

Fang, Cheng, et al. "Prospective validation of machine learning algorithms for absorption, distribution, metabolism, and excretion prediction: An industrial perspective." *Journal of Chemical Information and Modeling* 63.11 (2023): 3263-3274.


# Which Property Best Separates Soluble vs Insoluble Molecules?




## The Gini Index Quantifies the “Purity” of a Split

---


$$\text{Gini Index} = 1 - \sum_{i=1}^n (P_i)^2$$




1.00



0.83




0.66



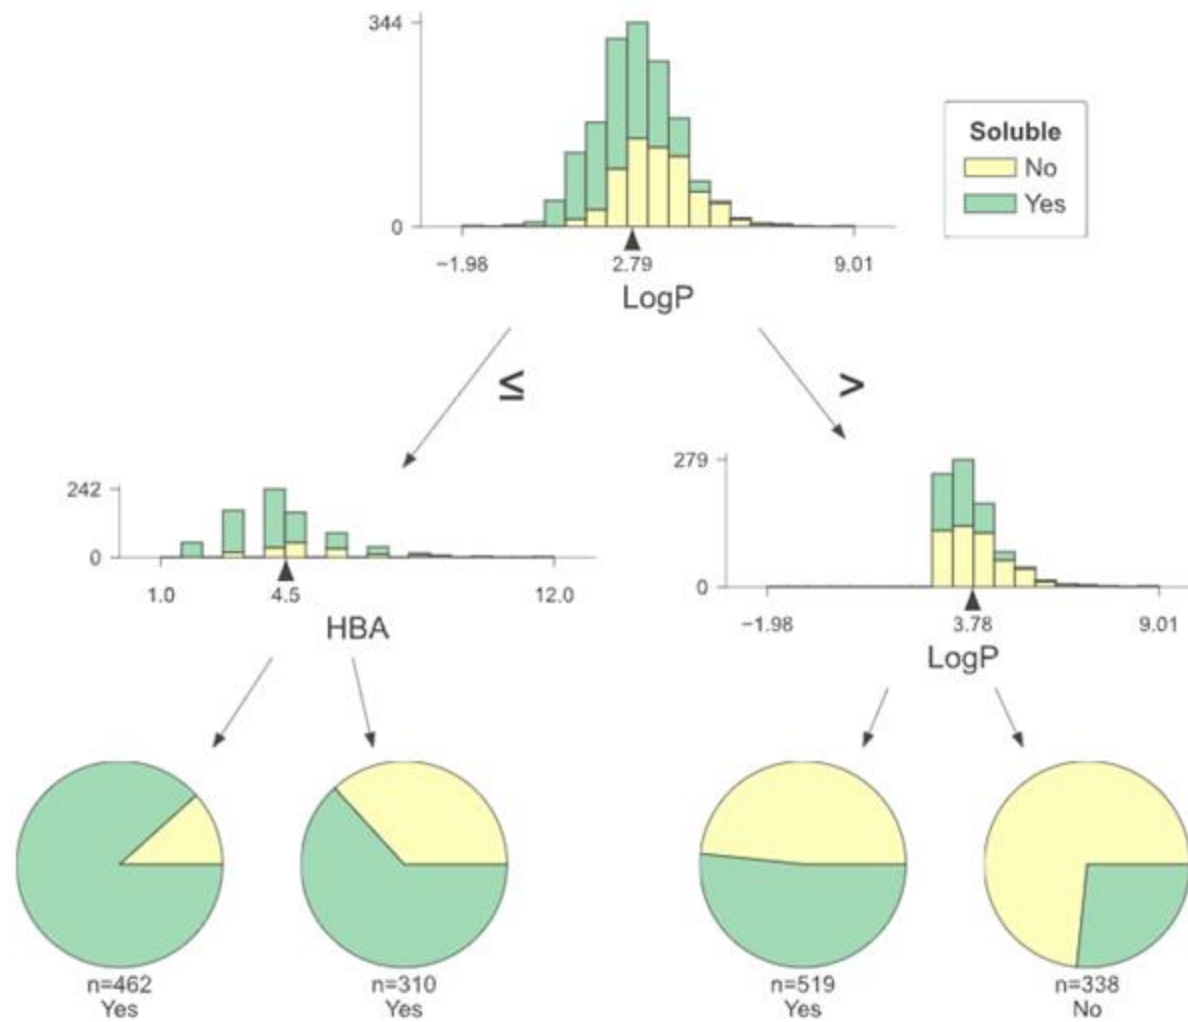
0.50

## Build a Decision Tree

---




# A Better (IMO) Decision Tree Visualization





# Random Forest Uses an Ensemble of Decision Trees

---



## There Are Many Tree Ensemble Methods

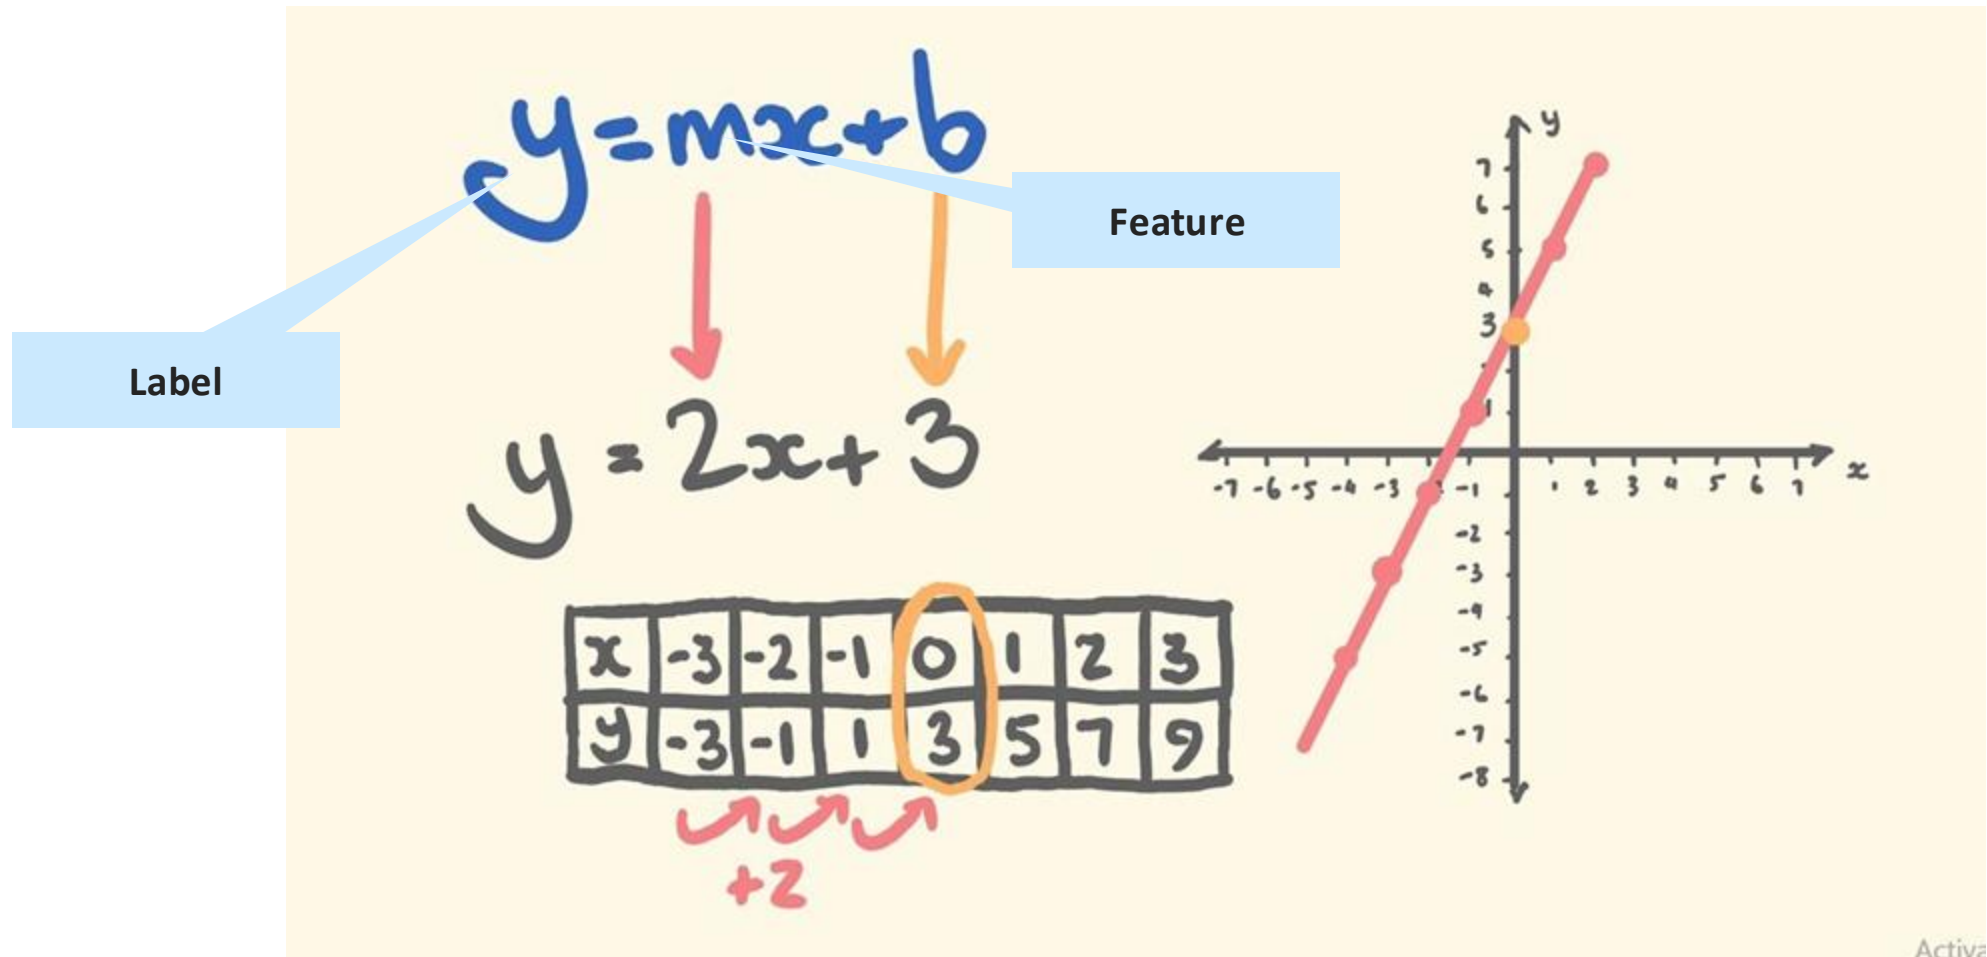
---




***XGBoost***



# ML Predicts a Set of Labels (y) Based on Features (X)





# Defining Labels (y) and Features (X) - An Example



$y = \text{Log Aqueous Solubility}$

Label



010110010001000111001000110010101010111

Features

$X = \text{A Vector Representing a Molecule}$



# Define Features Based on Books People Have Read



Fred	1	0	0	1	1	1	0	0	1	0	0	1	1	1	0	0
Sally	1	0	1	0	1	1	0	0	0	0	1	1	0	1	0	0
Jane	0	0	1	0	0	0	1	1	0	0	1	1	1	0	0	0

Create vectors representing books purchased by individuals  
1 = bought book  
0 = did not buy book


# Chemical Fingerprints as Molecular Descriptors



1	1	1	0	1	1	0	1	0
2	1	1	0	1	0	0	0	0




# Chemical Fingerprints as Vector Representations of Molecules



# There are A LOT of ways to do this

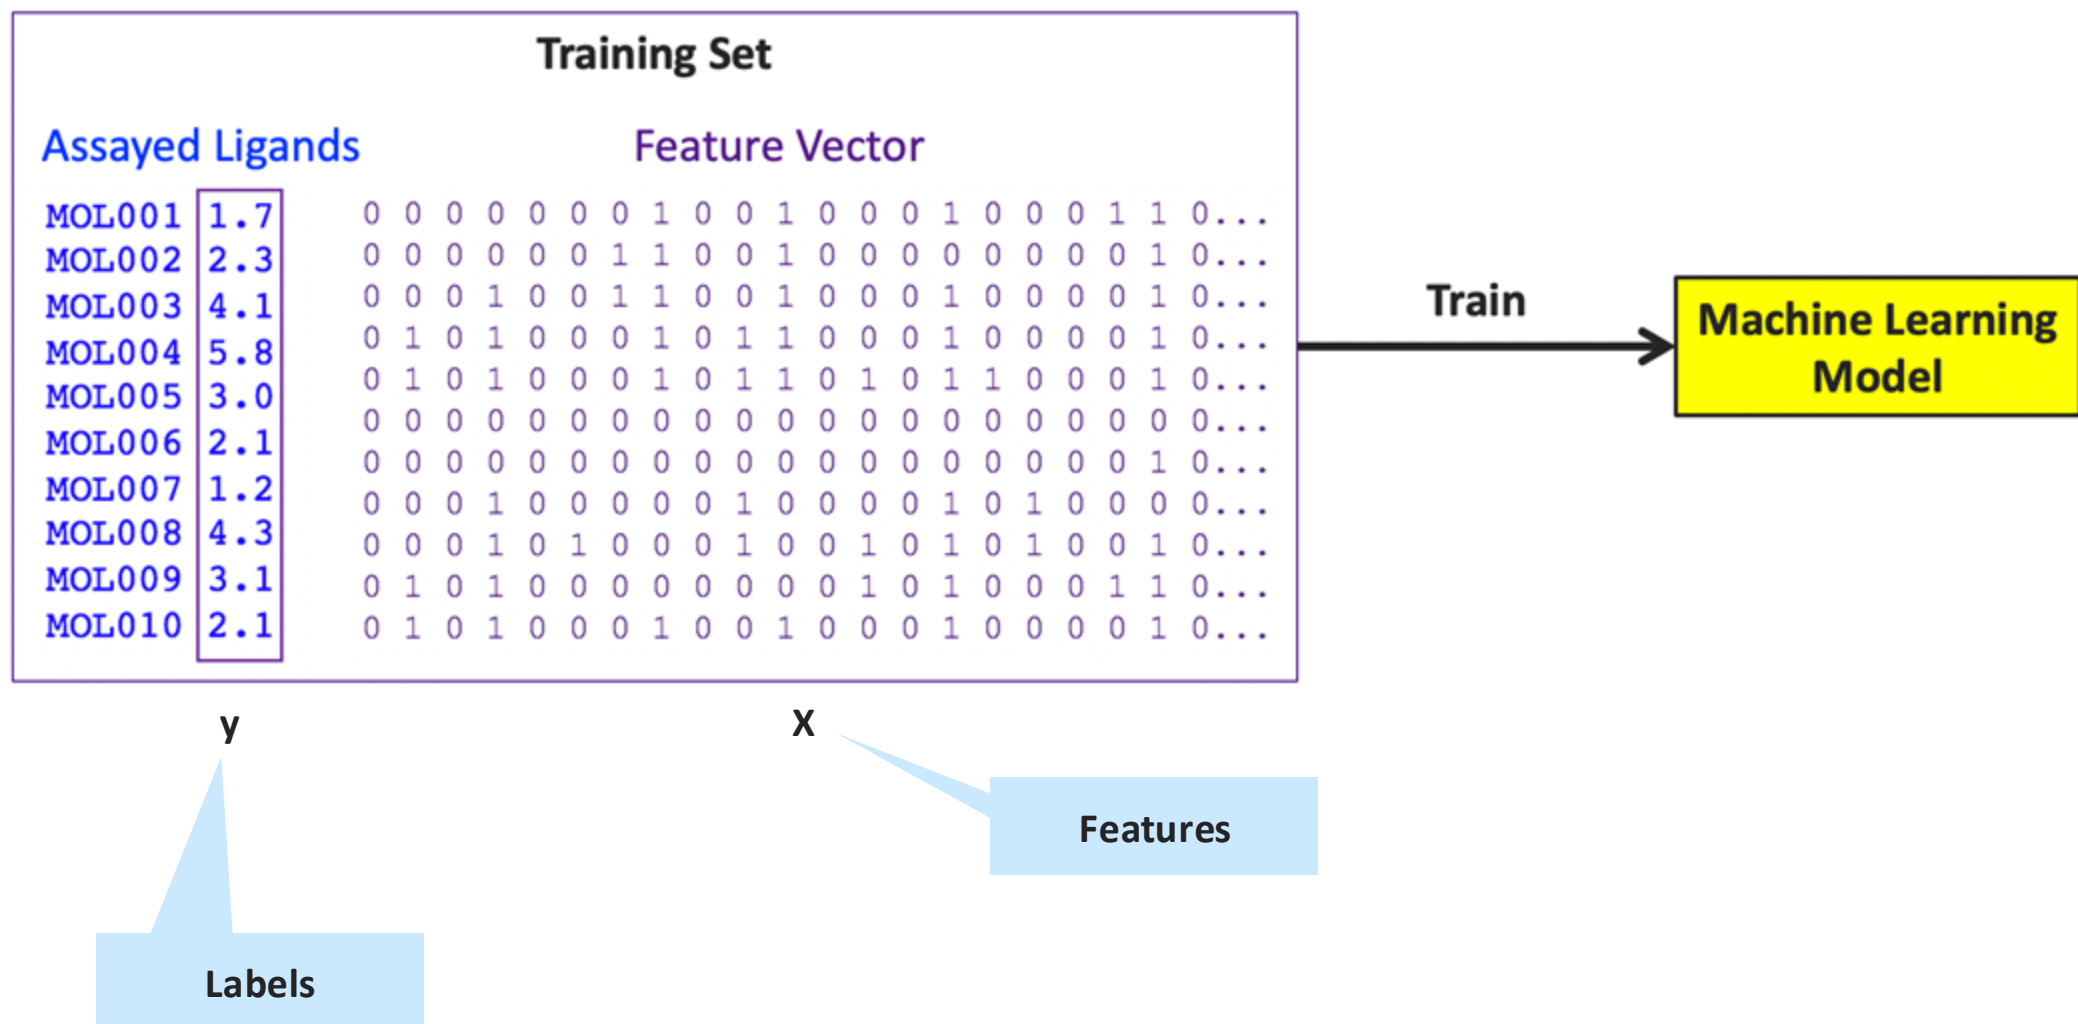
---



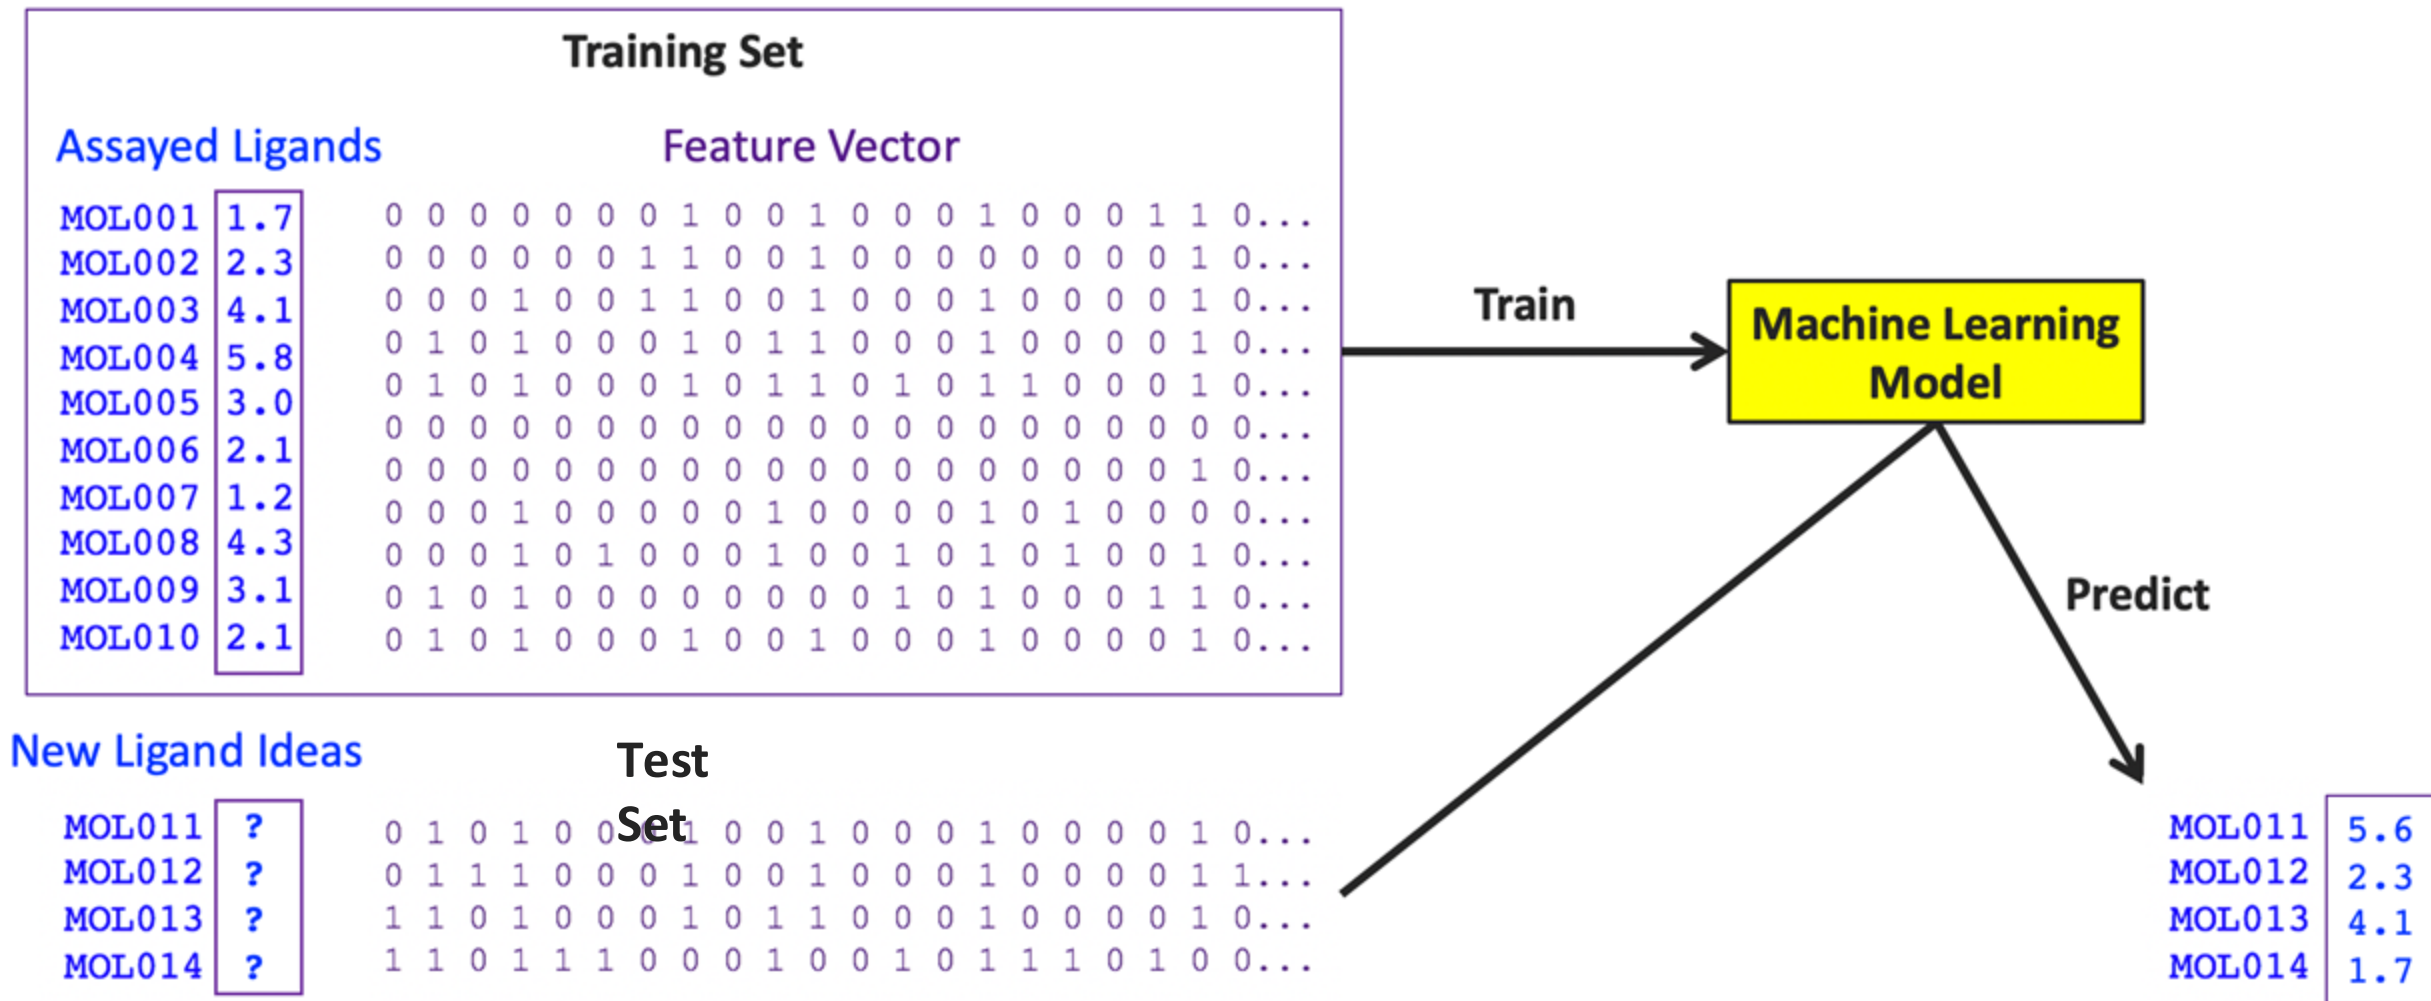
**2 Volumes**  
**6,000 references from 450 journals**

**DRAGON 7 has 5,270 descriptors z**


# Training a Machine Learning Model



# Making Predictions With a Machine Learning Model





# Prediction Performance of an Aqueous Solubility Model



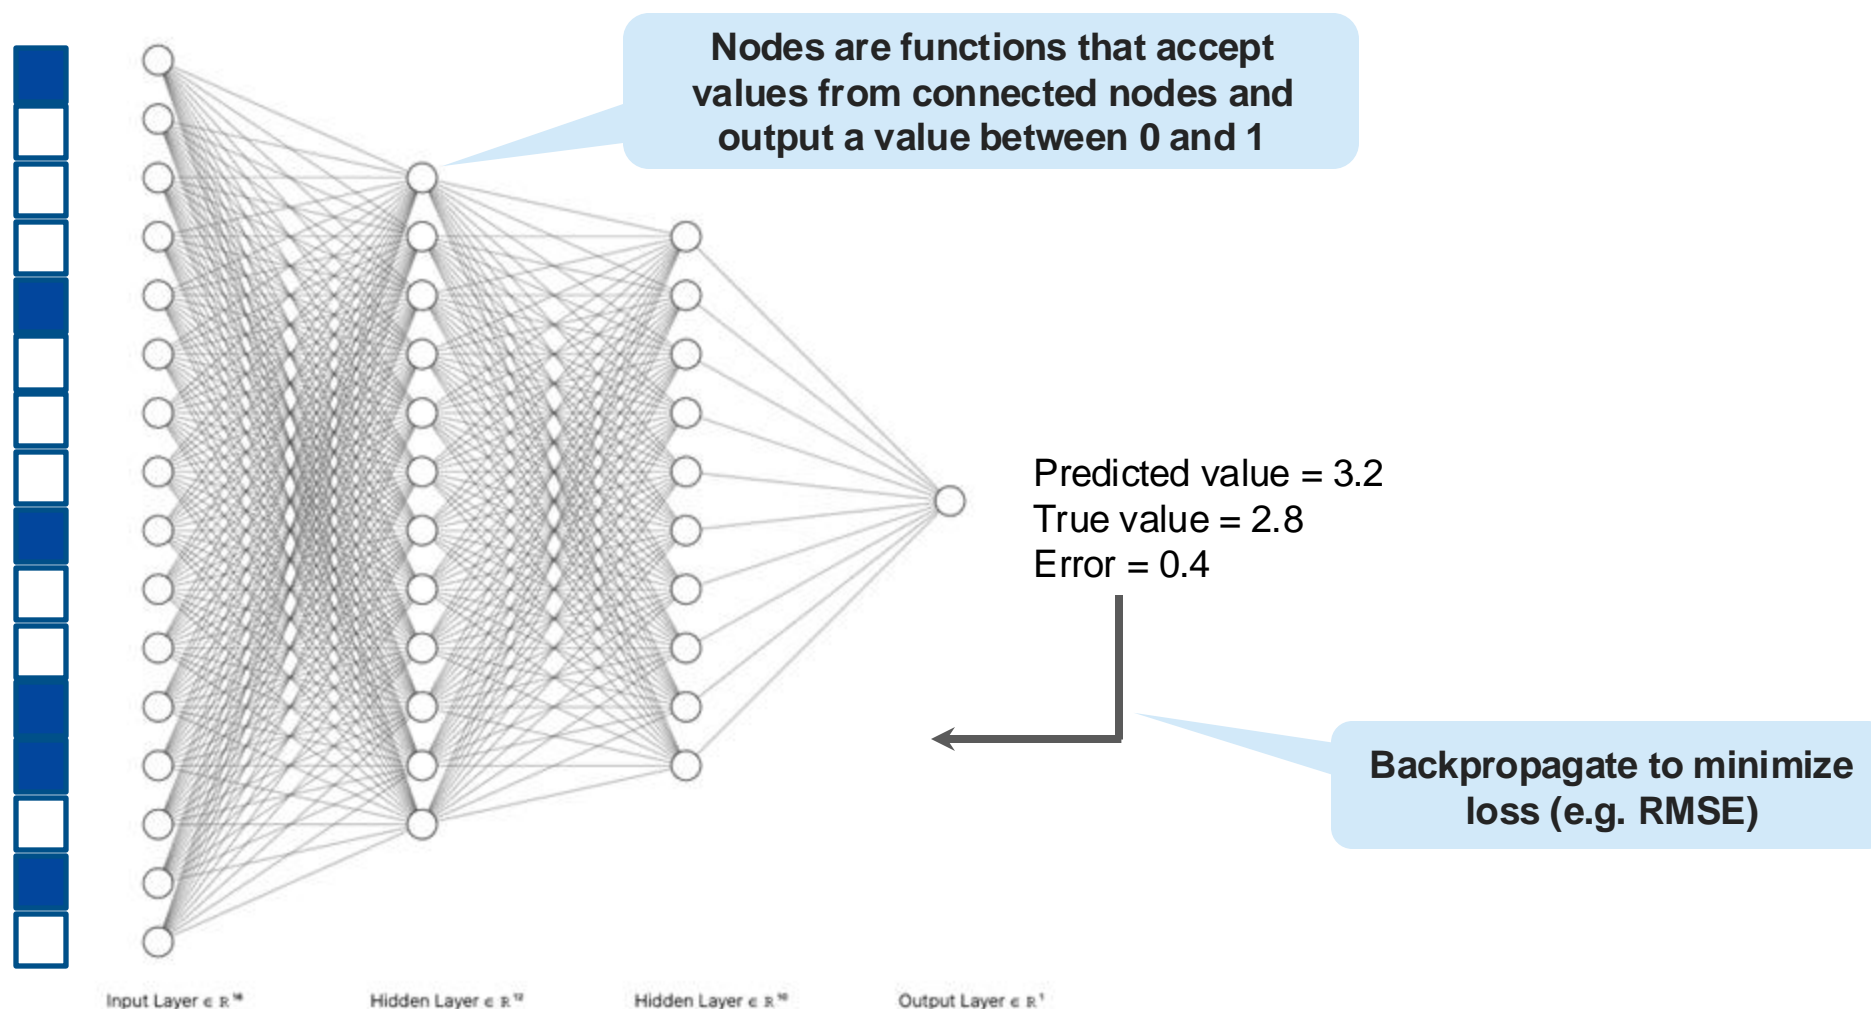
Dynamic range is very large,  
does not represent a typical use  
case

# Chemical Fingerprints as Vector Representations of Molecules





# Neural Networks Adjust Weights to Minimize a Loss Function




# Using Neural Networks to Create New Molecular Representations



## Graph Convolutions

*J. Comput. Aided Mol. Des.* 2016, 595–608



## Message Passing Neural Network

*J. Chem. Inf. Model.* 2019, 59, 3370–3388


# Are Neural Network Representations Better?

Dataset Size


<1000

1000-2000

>2000



# Incorporating 3D into Molecular Machine Learning is an Unsolved Problem



2D single-instance

→ 0 0 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0 1 1 0...


→ 5.21



3D single-instance

→ 0 0 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0 1 1 0...

→ 5.21



3D multiple-instance

0 0 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0 1 1 0...  
0 0 0 0 0 0 1 1 0 0 1 0 0 0 0 0 0 0 0 1 0...  
0 0 0 1 0 0 1 1 0 0 1 0 0 0 1 0 0 0 0 1 0...  
0 1 0 1 0 0 0 1 0 1 1 0 0 0 1 0 0 0 0 1 0...  
0 1 0 1 0 0 0 1 0 1 1 0 1 0 1 1 0 0 0 1 0...  
0 0...  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0...  
0 0 0 1 0 0 0 0 0 1 0 0 0 0 1 0 1 0 0 0 0...  
0 0 0 1 0 1 0 0 0 1 0 0 1 0 1 0 1 0 0 1 0...  
0 1 0 1 0 0 0 0 0 0 0 0 1 0 1 0 0 0 1 1 0...  
0 1 0 1 0 0 0 1 0 0 1 0 0 0 1 0 0 0 0 1 0...

→ 5.21



# How to Find Me

---

<https://pwalters.github.io>

My Boring Website

Publications

Tutorials

Blog

Videos



**Pat Walters**

Cheminformatics, ML

📍 Cambridge, MA

✉ Email

🔍 Google Scholar

🐙 Github

🌐 LinkedIn

✂ X (formerly Twitter)

Pat Walters is Chief Data Officer at Relay Therapeutics in Cambridge, MA. Prior to joining Relay, he spent more than 20 years at Vertex Pharmaceuticals where he was Global Head of Modeling & Informatics. Pat is the 2023 recipient of the [Herman Skolnik Award](#) for Chemical Information Science from the American Chemical Society. He is a member of the editorial advisory boards for the Journal of Chemical Information and Modeling and Artificial Intelligence in the Life Sciences, and previously held a similar role with the Journal of Medicinal Chemistry. Pat is co-author of the book "[Deep Learning for the Life Sciences](#)", published in 2019 by O'Reilly and Associates. He received his Ph.D. in Organic Chemistry from the University of Arizona where he studied the application of artificial intelligence in conformational analysis. Prior to obtaining his Ph.D., Pat worked at Varian Instruments as both a chemist and a software developer. He received his B.S. in Chemistry from the University of California, Santa Barbara.